

CS 161 Fall 2017: Section 6 Solutions

Currency Exchange

Suppose the various economies of the world use a set of currencies C_1, \dots, C_n —think of these as dollars, pounds, bitcoins, etc. Your bank allows you to trade each currency C_i for any other currency C_j , and finds some way to charge you for this service (in a manner to be elaborated in the subparts below). We will devise algorithms to trade currencies to maximize the amount we end up with.

- (a) Suppose that for each ordered pair of currencies (C_i, C_j) , the bank charges a flat fee of $f_{ij} > 0$ dollars to exchange C_i for C_j (regardless of the quantity of currency being exchanged). Devise an efficient algorithm which, given a starting currency C_s , a target currency C_t , and a list of fees f_{ij} for all $i, j \in \{1, \dots, n\}$, computes the cheapest way (that is, incurring the least in fees) to exchange all of our currency in C_s into currency C_t . Justify the correctness of your algorithm and its runtime.

Build the complete graph on the currencies with weights equal to the corresponding fees; i.e. $G = (V, E, w)$ where $V = \{C_1, \dots, C_n\}$, $E = \{(C_i, C_j) : i \neq j\}$, and $w(C_i, C_j) = f_{ij}$. Run Dijkstra's algorithm from C_s ; return a shortest $C_s \rightarrow C_t$ path.

Correctness: Notice that any path from C_s to C_t in G indicates a sequence of exchanges, and that its total weight is precisely the sum of the fees needed to perform those exchanges. Thus it suffices to find a $C_s \rightarrow C_t$ path in G of minimum weight. Note that the f_{ij} 's are all positive, so this is a valid input to Dijkstra's algorithm.

Running time: $O(n^2)$ total. Since all exchange pairs are possible, we have $m = \binom{n}{2} = \Theta(n^2)$ edges—this is also how long it takes to build G . Dijkstra's algorithm therefore takes $O(n^2 + n \log n) = O(n^2)$ time using the Fibonacci heap implementation. Other implementations are possible; for example, using a red-black tree or binary heap leads to $O((m + n) \log n) = O(n^2 \log n)$.

- (b) Consider the more realistic setting where the bank does not charge flat fees, but instead uses exchange rates. In particular, for each ordered pair (C_i, C_j) , the bank lets you trade one unit of C_i for $r_{ij} > 0$ units of C_j . Devise an efficient algorithm which, given starting currency C_s , target currency C_t , and a list of rates r_{ij} , computes a sequence of exchanges that results in the greatest amount of C_t . Justify the correctness of your algorithm and its runtime. [Hint: How can you turn a product of terms into a sum?]

Build $G = (V, E, w)$ as in part (a), but with weights $w(C_i, C_j) = -\log(r_{ij})$. Run Bellman-Ford from C_s and return a shortest $C_s \rightarrow C_t$ path.

Correctness: As in part (a), a sequence of exchanges corresponds to a path in G . However, here we want a path $C_{i_1} = C_s, C_{i_2}, \dots, C_{i_k} = C_t$ that maximizes $\prod_{\ell=1}^{k-1} r_{i_\ell, i_{\ell+1}}$. Since \log is a monotonically increasing function (i.e. if $a \geq b$, then $\log(a) \geq \log(b)$), this is the same as maximizing $\log\left(\prod_{\ell=1}^{k-1} r_{i_\ell, i_{\ell+1}}\right) = \sum_{\ell=1}^{k-1} \log(r_{i_\ell, i_{\ell+1}})$. Finally, this is equivalent to minimizing $\sum_{\ell=1}^{k-1} -\log(r_{i_\ell, i_{\ell+1}}) = \sum_{\ell=1}^{k-1} w(C_{i_\ell}, C_{i_{\ell+1}})$, which is the shortest path objective. Note that we must use Bellman-Ford rather than Dijkstra's algorithm, since these weights may be negative.

Running time: $O(n^3)$ total. G can be built in time $O(n^2)$ time, and Bellman-Ford takes $O(n^3)$ time since we have $\Theta(n^2)$ edges.

- (c) Due to fluctuations in the markets, it is occasionally possible to find a sequence of exchanges that lets you start with currency A, change into currencies B, C, D, etc., and then end up changing back to

currency A in such a way that you end up with more money than you started with—that is, there are currencies C_{i_1}, \dots, C_{i_k} such that

$$r_{i_1 i_2} \times r_{i_2 i_3} \times \dots \times r_{i_{k-1} i_k} \times r_{i_k i_1} > 1.$$

Devise an efficient algorithm that finds such an anomaly if one exists. Justify the correctness of your algorithm and its runtime.

Run Bellman-Ford on the same graph as in part (b); then execute one more iteration of Bellman-Ford to check if there is a negative cycle in G . If there is, the cycle is the anomaly—trading currencies along the cycle will result in a profit.

Correctness: A currency anomaly $\prod_{\ell=1}^{k-1} r_{i_\ell, i_{\ell+1}} > 1$ implies (by the same log manipulations we did in part (b)) that $\sum_{\ell=1}^{k-1} w(C_{i_\ell}, C_{i_{\ell+1}}) = \sum_{\ell=1}^{k-1} -\log(r_{i_\ell, i_{\ell+1}}) < 0$. Thus there is a negative cycle in G , which can be found by an extra iteration of Bellman-Ford.

Running time: $O(n^3)$ total. We are doing the same thing as in part (b), plus one extra iteration which takes $O(|E|) = O(n^2)$ time.

Traveling Across the Country

We have a graph representation of the country, where nodes u_i are on the east coast and nodes v_j are on the west coast, with n nodes in total. We also have $|E|$ undirected edges representing distances between these cities.

- (a) Design an efficient algorithm to compute the shortest path starting at *any* city on the east coast and ending at *any* city on the west coast.

Add two nodes to the graph: one node E connected to all cities on the east coast with edge weight 0, and one node W connected to all cities on the west coast also with edge weight 0. Run Dijkstra's on the new graph starting at E , and return the path/length to W . The asymptotic runtime is the same as that of Dijkstra's on the same graph.

- (b) This time, we start from a specific city u_i and end at a specific city v_j . However, we impose an additional restriction that we must traverse one of the edges between two cities 3 times: that is, for some w, w' , we must traverse $w \rightarrow w'$, $w' \rightarrow w$, and then $w \rightarrow w'$ again. Design an efficient algorithm to find the shortest path from u_i to v_j with this additional constraint.

Run Dijkstra's once starting at u_i and once starting at v_j . Now, we have the distance/path from every city to u_i and v_j . For each edge (a, b) in the graph, check the path length given by $D(u_i, a) + 3W_{ab} + D(v_j, b)$ and $D(u_i, b) + 3W_{ab} + D(v_j, a)$. Take the minimum length path. The asymptotic runtime is the same as that of Dijkstra's because we run Dijkstra's twice and iterate over all edges once.